Lab 6

Thomas Haase

October 10, 2025

Table of contents

Part 1: Meta-learners for job training evaluation 1
- Task 1 2
- Task 2 2
- Task 3 3
- Task 4 3
- Task 5 5

Part 2: Heterogeneity I 6
- Task 1 6
- Task 2 7
- Task 3 7
- Task 4 7
- Task 5 9

Part 1: Meta-learners for job training evaluation

The dataset "job_training_updated.csv" contains information about 12,000 individuals who either participated or did not participate in a job training program, including: - training: Binary indicator of whether individual participated in training (treatment) - earnings: Post-training annual earnings in thousands of dollars (outcome) - age: Age of individual - education: Years of education - prior_earnings: Earnings before training program - employment_history: Years of prior employment - urban: Binary indicator of urban residence

```
library(easystats)
library(data.table)
library(kableExtra)

library(rpart)
library(rpart.plot)


library(randomForest)

library(caret)

library(htetree)
library(grf)

set.seed(5)

setwd("~/Github/ML-Labs/6")
```

Task 1

1. Fit regular OLS regression using lm(), including all non-treatment and non-outcome variables as control variables. Interpret the coefficient for the treatment variable as the average treatment effect. Considering what we talked about in the lecture, what properties of the data would lead you to believe your estimate is biased? Motivate.

```
d <- data_read("job_training_updated.csv")

m1 <- lm(earnings ~ training + age + education +
         prior_earnings + employment_history + urban, d)

m1 |> report_table() |> summary() |> print_md()
```

Parameter	Coefficient	95% CI	t(11993)	p	Std. Coef.	Fit
(Intercept)	2.35	(-1.81, 6.51)	1.11	0.269	−1.49e-16	
training	24.86	(23.83, 25.89)	47.24	< .001	0.40	
age	−0.19	(-0.34, −0.03)	−2.34	0.019	−0.05	
education	5.86	(5.63, 6.09)	49.37	< .001	0.41	
prior earnings	−0.14	(-0.15, −0.13)	−28.37	< .001	−0.25	
employment history	0.79	(0.64, 0.94)	10.34	< .001	0.24	
urban	2.57	(1.68, 3.47)	5.63	< .001	0.04	

AICc	1.11e+05
R2	0.39
R2 (adj.)	0.39
Sigma	24.40

The ATE of the treatment variable `training` on earnings is very large. Despite the statistical significance the model could be improved by causal modeling. The standard paradigm is problematic because it lacks assumes linearity. In case a confounder is non-linear the estimate will be biased.

Task 2

Next, you shall estimate an orthogonal learner, using decision trees as the method for predicting both the treatment and the outcome. Please follow the following steps:

- a. Train a decision tree model using rpart() to predict training from all confounders using the full dataset. For classification trees, use method="class" and for the control parameters use: cp=0, minbucket=5, maxdepth=30 (i.e., control=rpart::rpart.control(cp=0,minbucket=5, maxdepth=30)).
- b. Train a decision tree model using rpart() to predict earnings from all confounders using the full dataset. For regression trees, use method="anova" and the same control parameters: cp=0, minbucket=5, maxdepth=30 (i.e., control=rpart::rpart.control(cp=0,minbucket=5, maxdepth=30)).
- c. Make predictions of treatment (using model from a, with type="prob") and outcome (using model from b) for all observations.
- d. Calculate residuals for all observations: X_tilde = X - X_hat, Y_tilde = Y - Y_hat.
- e. Estimate the ATE by regressing Y_tilde on X_tilde using lm().
- f. Report the ATE. How does it compare to your OLS estimate in #1?

- g. Which of the two methods do you trust more? Can you think of any aspect of the implementation of the orthogonal learner which could bias its estimate?

```
# a
m_X <- rpart(training ~ age + education + prior_earnings +
               employment_history + urban, d, method = "class",
             control = rpart.control(cp = 0, minbucket = 5,
                                     maxdepth = 30))
# b
m_Y <- rpart(earnings ~ age + education + prior_earnings +
               employment_history + urban, d, method = "anova",
             control = rpart.control(cp = 0, minbucket = 5,
                                     maxdepth = 30))
# c
X_hat <- predict(m_X, newdata = d[3:7], type="prob")[,"1"]
Y_hat <- predict(m_Y, newdata = d[3:7])

# d
residuals <- data.frame(X = d$training - X_hat,
                        Y = d$earnings - Y_hat)

# e
m_ate <- lm(Y ~ X, residuals)

# f
m_ate |> parameters() |> print_md()
```

Parameter	Coefficient	SE	95% CI	t(11998)	p
(Intercept)	−2.54e-15	0.10	(-0.20, 0.20)	−2.46e-14	> .999
X	9.72	0.31	(9.12, 10.33)	31.40	< .001

```
m_ate |> report_parameters(include_intercept = F)
```

```
  - The effect of X is statistically significant and positive (beta = 9.72, 95% CI [9.12,
10.33], t(11998) = 31.40, p < .001; Std. beta = 0.28, 95% CI [0.26, 0.29])
```

Task 3

Given your conclusions in 2, do you think either of the following two changes to the setup of the orthogonal learner could improve the ATE estimate? (i) switching from a decision tree to a random forest, (ii) add cross-fitting. Motivate.

Cross-fitting means to predict out-of-fold to block contamination whereby the orthogonal learner does not bias the residuals in hold-out towards 0. To adress the high variance of the trees by prohibiting the most important variables to dominate the trees. The negative aspect of forests that they are less interpretable is not relevant here since we are only interested in predicting the effect of the confounders to X and Y.

Task 4

Now you shall implement the two updates discussed in 3. Please do the following:

a. Divide your data into 5 folds (hint: you can use createFolds() from the caret package).
b. Create a for-loop which in each iteration i does the following:

i. Train a random forest model using randomForest() (with ntree=200 and mtry=2) predicting training from confounders on data in folds = i.
ii. Train a random forest model using randomForest() (with ntree=200 and mtry=2) predicting earnings from confounders on data in folds = i.
iii. Use models from (i) and (ii) to predict treatment (with type="prob") and outcome for ob- servations in fold i.
iv. Calculate residuals X_tilde and Y_tilde for observations in fold i.
v. Store residuals from fold i.

c. Combine dataset of residuals and regress Y_tilde on X_tilde using lm().
d. Report the estimated ATE. Do you trust this estimate more than those in 2, and if so why (or why not)?

```
if(file.exists("Task_1_4.rds")){
  residuals <- readRDS("Task_1_4.rds")
} else{

  # a
  ids_folds <- createFolds(d$earnings, k = 5)
  residuals_list <- list()

  # b
  for(i in seq_along(ids_folds)){

    ids_test <- ids_folds[[i]]

    cat("------------------------------\n
        Start Iteration", i, "\nSplitting Data\n")

    testdata <- d[ids_test,] |>
      as.data.frame() |>
      data_select(c("age","education","prior_earnings",
                    "employment_history","urban",
                    "training","earnings"))

    trainingdata <- d[-ids_test,] |>
      as.data.frame() |>
      data_select(c("age","education","prior_earnings",
                    "employment_history","urban",
                    "training","earnings"))

    confounders <- c("age","education","prior_earnings",
                    "employment_history","urban")

    # i
    cat("Calculate RF for X\n")
    m_X <- randomForest(x = trainingdata[,confounders],
                        y = trainingdata$training,
                        ntree = 200, mtry = 2)
    # ii
    cat("Calculate RF for Y\n")
    m_Y <- randomForest(x = trainingdata[,confounders],
                        y = trainingdata$earnings,
                        ntree = 200, mtry = 2)
    # iii
    cat("Predict Testdata\n")
```

```
    X_hat <- predict(m_X, newdata = testdata[,confounders])
    Y_hat <- predict(m_Y, newdata = testdata[,confounders])

    # iv
    cat("Store Residuals\n")
    residuals_list[[i]] <- data.frame(
      fold = i,
      original_index = ids_test,
      X = testdata$training - X_hat,
      Y = testdata$earnings - Y_hat
    )
  }

  residuals <- do.call(rbind, residuals_list)
  residuals <- residuals[order(residuals$original_index), ]

  saveRDS(residuals,"Task_1_4.rds")
}


# c
m_ate <- lm(Y ~ X, residuals)

# f
m_ate |> parameters() |> print_md()
```

Parameter	Coefficient	SE	95% CI	t(11998)	p
(Intercept)	−0.20	0.14	(-0.47, 0.07)	−1.44	0.151
X	12.94	0.35	(12.26, 13.62)	37.37	< .001

```
m_ate |> report_parameters(include_intercept = F)
```

```
  - The effect of X is statistically significant and positive (beta = 12.94, 95% CI [12.26,
13.62], t(11998) = 37.37, p < .001; Std. beta = 0.32, 95% CI [0.31, 0.34])
```

I trust this estimate more then the ones before since random forest prevents overfitting the ATE is not biased towards 0 anymore. Because of that the RF ATE is a bit larger then the decisiontree ATE.

Task 5

Suppose we learn that the true average treatment effect is 5.5 thousand dollars. Report which method came closest, and discuss what this says about the properties of the data—in particular the relation between the confounders and the treatment and outcome.

The used models assume all confounding variables are included in the model. Since the true ATE is so different from the estimated ones not all confounding variables were included in the model. This points the researcher towards theorybuilding :)

Part 2: Heterogeneity I

The dataset "scholarship.csv" contains information about 15,000 students who either received or did not receive a college scholarship, including: - scholarship: Binary indicator of scholarship receipt (treatment) - completed: Binary indicator of degree completion within 6 years (outcome) - gpa: High school GPA (scale 0-4) - parental_income: Parental income in thousands of dollars - first_generation: Binary indicator of first-generation college student status - sat_score: SAT score (scale 400-1600) - distance_to_college: Distance from home to college in miles - financial_need: Measure of financial need (scale 0-100)

```
rm(list = ls())

d <- data_read("scholarship.csv")
```

Task 1

Suppose your co-author, who has done a careful literature review, has found support for two of the variables, first_generation and financial_need, having a moderating effect. What you shall do first is examine whether you find evidence of this in your data. Implement a standard linear regression using lm() (or glm() if you prefer logistic regression) with the treatment variable as well as all other input variables (presumed confounders) included, with first_generation and financial_need interacted with the treatment variable. Report your findings: do you find evidence supporting your colleague's conclusion from the literature?

```
m1 <- lm(completed ~ scholarship + gpa + parental_income +
         first_generation + sat_score + distance_to_college + financial_need +
         scholarship:first_generation + scholarship:financial_need,
         data = d)

m1 |> report_table() |> summary() |> print_md()
```

Parameter	Coefficient	95% CI	t(14990)	p	Std. Coef.	Fit
(Intercept)	−0.16	(-0.23, −0.08)	−4.19	< .001	−0.04	
scholarship	−0.05	(-0.09, −9.86e-03)	−2.46	0.014	0.13	
gpa	0.20	(0.19, 0.21)	29.98	< .001	0.23	
parental income	5.28e-04	(3.85e-04, 6.72e-04)	7.22	< .001	0.08	
first generation	−0.01	(-0.04, 0.01)	−0.99	0.321	0.11	
sat score	2.43e-04	(2.01e-04, 2.85e-04)	11.32	< .001	0.09	
distance to college	−9.03e-04	(-1.24e-03, −5.69e-04)	−5.29	< .001	−0.04	
financial need	−7.19e-04	(-1.43e-03, −1.11e-05)	−1.99	0.047	0.02	
scholarship × first generation	0.19	(0.16, 0.22)	12.09	< .001	0.11	
scholarship × financial need	1.87e-03	(1.08e-03, 2.66e-03)	4.65	< .001	0.04	
AICc						14718.71
R2						0.11
R2 (adj.)						0.11

Parameter	Coefficient	95% CI	t(14990)	p	Std. Coef.	Fit
Sigma						0.40

```
m1 |> parameters() |> plot(show_intercept = T)
```

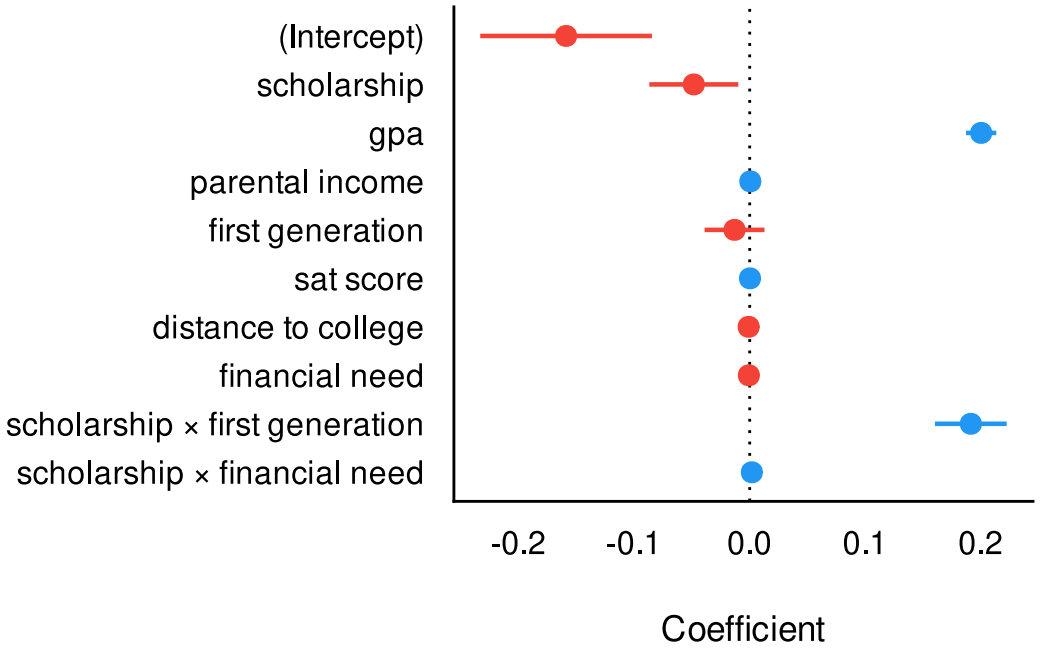

My co-author is partly right: - “first generation” has a statistically significant positive interaction effect - “financial need” has a statistically significant positive interaction effect, but it is so tiny that the significance is probably due to the large observation size. It does not seem to exist.

Task 2

Considering what we discussed in the lecture, what is one limitation of this standard approach to effect heterogeneity? What are properties of the data (or state of the field) that could make this limitation more or less problematic?

Research frequently find that the effect of events, exposures, policies (etc) vary across individuals. Usually interactionterms are implemented in the model with the selection of subgroups being based on theory or convention. Limitations of this procedure are that it assumes we know moderators beforehand. While exploratory analysis could help it creates risk of p-hacking and is not feasable for data with too many variables.

Task 3

Next, you shall consider an alternative approach to effect heterogeneity, using causal trees. At a high level, describe what is the key difference in assumption we make when using causal trees compared to the traditional approach?

Causal trees identify the splits that maximize across-leaf variation in the within-leaf treated–untreated outcome difference. Standard trees just minimize the variance in the leaves. Using trees solves the problem above since they include the most relevant interactions by design.

Task 4

Perform a causal tree analysis by doing the following:

a. Estimate the causal tree using the function causalTree() from the htetree package, specifying the formula as in #1 (except drop the interactions and leave out the treatment variable; the latter is specified separately). Use the following parameters: split.Rule="CT", cv.option="CT", split.Honest=TRUE, split.Bucket=TRUE, minsize=60, cp=0, bucketNum=40.

b. Visualize the tree using rpart.plot() and describe the combination of splits which identify the population with (a) the largest treatment effect and (b) the smallest.

c. Suppose our dataset is a standard observational dataset common to the social sciences, e.g., a survey dataset of a random sample of the population. Given this information, what could be a potential threat to the validity of our causal tree results?

```
# a.
set.seed(5)
if(file.exists("task2_4.rds")){
  m2 <- readRDS("task2_4.rds")
  } else {

    m2 <- causalTree(
      completed ~ gpa + parental_income +
             first_generation + sat_score + distance_to_college + financial_need,
      data      = d,
      treatment = d$scholarship,
      split.Rule = "CT", cv.option = "CT",
      split.Honest = T, split.Bucket = T,
      minsize = 60, cp = 0.00001, bucketNum = 40
    )

    saveRDS(m2, "task2_4.rds")
}

# b
m2 |> rpart.plot()
```

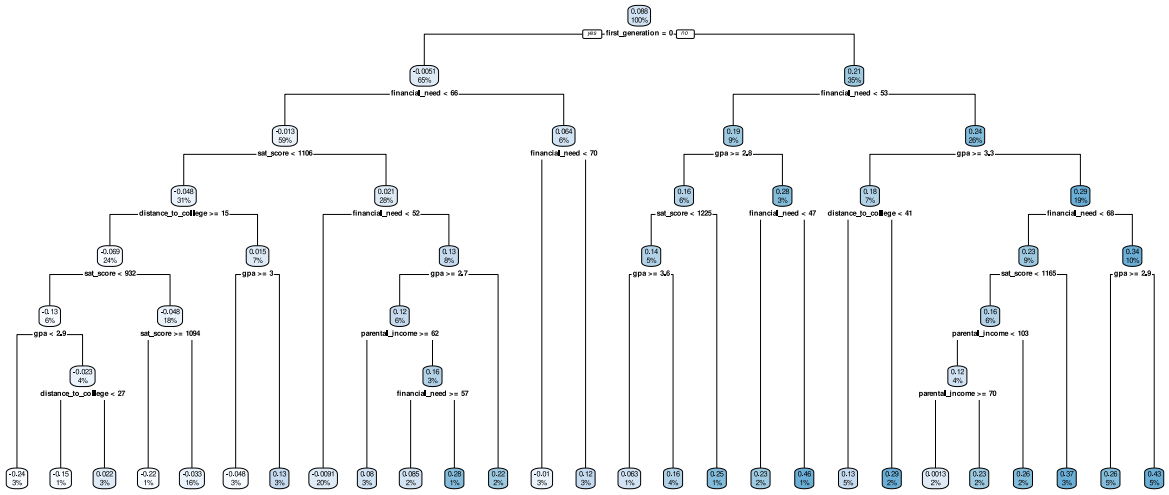

The population with the biggest treatment effect are students with - no first generation student - have a financial need of at least 68 - have a gpa below 2.9

The population with the smallest treatment effect are students with - the first generation - sat score of less than 932 - distance to college of more than 15 miles - gpa below 2.9

c. We included all variables, so we suppose all variables have an influence on the outcome, but that is not necessarily so.

Task 5

Given potential concerns of selection bias, you shall next examine the potential imbalance of treated and untreated observations inside different leaves. To do so, please follow these steps (Hint: various code-chunks are provided that may be helpful):

a. Estimate a propensity score model—a standard logistic regression model using glm() with family=binomial()—predicting the treatment variable based on the confounders. Specify type="response" in the predict() function.